



Scikit-Learn Framework

INTRODUCTION TO
SCIKIT-LEARN

WHAT IS SCIKIT-LEARN

- Scikit-Learn is one of the best-known Python packages for machine learning
 - scikit-learn.org
 - It provides efficient versions of many common ML algorithms
 - It has a clean, uniform, and streamlined API
 - Once you understand the syntax of Scikit-Learn, it is straightforward to use any model inside the library
 - In Python, Scikit-Learn is called with `sklearn`
 - **`import sklearn`**
-

OBJECT-ORIENTED MODELING

- A model is an **object**:
 - It has **state** (learned parameters)
 - It has **methods** (fit, predict)
- This allows:
 - Reuse
 - Inspection
 - Composition (pipelines)

```
from sklearn.linear_model import  
LinearRegression
```

```
# Model instantiation  
model = LinearRegression()
```

```
#Model fit  
model.fit(X, y)  
model.predict(X_new)
```

```
#Model learned parameters:  
coefficients = model.coef_
```

ESTIMATORS

- **Estimators:**

- Any object that can estimate a parameter based on a dataset
- The `.fit()` method performs the estimation

- **Two Forms:**

- Transformer
- Predictor

TRANSFORMERS

- **Transformers:**

- Estimators that can also transform a dataset
- Transformation is performed with the `.transform()` method
- The `.fit_transform()` method fits and transforms in one step
- Example: `StandardScaler`

```
import numpy as np
from sklearn.preprocessing import StandardScaler

# Feature matrix X (2D)
X = np.array([[1.0, 10.0],
              [2.0, 20.0],
              [3.0, 30.0]])

# 1. Instantiate
scaler = StandardScaler()

# 2. Compute mean/std and scale the data (mean = 0, std = 1)
X_scaled = scaler.fit_transform(X)

print(X_scaled)
# Output:
# [[-1.22474487 -1.22474487]
#  [ 0.          0.         ]
#  [ 1.22474487  1.22474487]]
```

PREDICTORS

- **Predictors:**
 - Estimators that are capable of making predictions
 - Has a `.predict()` method
 - Usually have a built-in `.score()` method
 - Example: `LinearRegression`

```
import numpy as np
from sklearn.linear_model import LinearRegression

# Feature matrix X (2D) and target vector y (1D)
X = np.array([[1], [2], [3], [4]])
y = np.array([2, 4, 6, 8])

# 1. Instantiate the estimator
model = LinearRegression()

# 2. Fit the estimator to the training data
model.fit(X, y)

# 3. Make predictions on new feature data
X_new = np.array([[5], [6]])
predictions = model.predict(X_new)

print("Predictions:", predictions) # Output: [10. 12.]
```

SIMPLE SCIKIT-LEARN USAGE

1. Prepare the data into a feature matrix and a target vector
2. Choose a model and import the estimator *class* from Scikit-Learn
3. Instantiate the model class and choose *hyperparameters* (not all model classes have hyperparameters)
4. Fit the model by calling the `.fit()` method
5. Apply the model with the `.predict()` or the `.transform()` method

Class: Scikit-Learn is built on Python classes and object-oriented principles

Hyperparameters: settings you choose before fitting; they control learning and are not learned from the data

PREPARING YOUR DATA

- The data should be in tabular form
 - Generally, NumPy or Pandas objects
 - The rows should be the observations or samples
- The feature matrix (X) should be a 2D NumPy array or a Pandas DataFrame
 - This applies even when there is only 1 feature
 - **Note:** if only one feature, it often must be reshaped using `.reshape(-1, 1)`
- The target (y) should be a 1D NumPy array or a Pandas Series

EXAMPLE

```
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split

# Create simple data
X = np.arange(1, 21).reshape(-1, 1)    # feature matrix (n_samples, n_features)
y = 2 * X.ravel()                      # target vector

# Hold out a test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)

# Fit on the training data only
model = LinearRegression()
model.fit(X_train, y_train)

# Predict on the unseen test data
y_pred = model.predict(X_test)
print("Slope:", model.coef_[0])
print("Intercept:", model.intercept_)
```

WHY USE SCIKIT-LEARN PIPELINES?

- **A pipeline treats the full workflow as one estimator:**
 - Prevents leakage during train/test splits and cross-validation
 - Applies the same preprocessing to training, validation, and new data
 - Keeps code concise and reproducible: one object for `.fit()`, `.predict()`, and `.score()`

HOW TO BUILD A PIPELINE: SCALE + KNN

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsRegressor

knn_pipe = Pipeline([
    ("scale", StandardScaler()),
    ("model", KNeighborsRegressor(n_neighbors=5)),
])

knn_pipe.fit(X_train, y_train)
y_pred = knn_pipe.predict(X_test)
```

DIFFERENT STEPS FOR DIFFERENT COLUMNS

```
numeric = Pipeline([
    ("impute", SimpleImputer(strategy="median")),
    ("scale", StandardScaler()),
])
preprocess = ColumnTransformer([
    ("num", numeric, ["age", "income"]),
    ("cat", OneHotEncoder(handle_unknown="ignore"), ["city"]),
])
model = Pipeline([
    ("prep", preprocess),
    ("reg", LinearRegression()),
])
model.fit(X_train, y_train)
```

Numerical columns

impute -> scale

Categorical columns

one-hot encode

Combined features

Linear Regression

REVIEW THE DOCUMENTATION!

- A large part of learning Scikit-Learn is getting comfortable with the documentation
- The documentation is vast and sometimes intimidating but learning how to read it will help you know how to implement the right model for your situation
- <https://scikit-learn.org/stable/>