




CLASSIFICATION METRICS



REGRESSION VS CLASSIFICATION (REVIEW)

- Regression and Classification are both supervised learning methods
- Regression
 - The target or response variable (y) is numeric
 - Model predicts the *value* of y
 - Prediction accuracy is measured by how close the predicted value is to the truth, such as

$$(y_i - \hat{y}_i)^2 \quad \text{or} \quad |y_i - \hat{y}_i|^2$$



ACCURACY AND ERROR RATE

- Classification predicts categorical outcomes or class probabilities.

$$\mathbf{ErrorRate} = \frac{1}{n} \sum_{i=1}^n 1\{y_i \neq \hat{y}_i\}$$

$$\mathbf{Accuracy} = \frac{\text{Correct predictions}}{\text{Total predictions}} = 1 - \text{Error rate}$$

- Both metrics treat all errors equally, regardless of their cost.
 - With class imbalance, predicting only the majority class can yield high accuracy.
-

CONFUSION MATRIX

		Predicted Output	
		Positive (PP)	Negative (PN)
True Values	Positive (P)	True Positives TP	False Negatives FN
	Negative (N)	False Positives FP	True Negatives TN



PRECISION

		Predicted Output	
		Positive (PP)	Negative (PN)
True Values	Positive (P)	True Positives TP	False Negatives FN
	Negative (N)	False Positives FP	True Negatives TN

- Of those that are **predicted to be positive**, the percent that are actually positive
- Describes how good the model is at predicting the positive class
- Also call **positive predictive value**

$$Precision = \frac{TP}{TP + FP}$$

Under what conditions is high **precision** the most appropriate performance objective for a classification model?



RECALL

		Predicted Output	
		Positive (PP)	Negative (PN)
True Values	Positive (P)	True Positives TP	False Negatives FN
	Negative (N)	False Positives FP	True Negatives TN

- Of those that are **actually positive**, the percent that are predicted to be positive
- The proportion of positives that are correctly identified
- Also called **sensitivity** or **true positive rate**

$$Recall = \frac{TP}{TP + FN}$$

Under what conditions is high **recall** the most appropriate performance objective for a classification model?

F1 SCORE

- The precision & recall are often combined into one metric called F_1
- F_1 is the **harmonic mean** of the precision and recall

$$F_1 = \frac{2}{\frac{1}{precision} + \frac{1}{recall}} = 2 \left(\frac{precision \cdot recall}{precision + recall} \right)$$

- The F_1 score favors classifiers with similar precision and recall scores
-

F_β SCORE

- The F_1 score can be generalized to give more weight to either the precision or recall

$$F_\beta = (1 + \beta^2) \left(\frac{\textit{precision} \cdot \textit{recall}}{(\beta^2 \cdot \textit{precision}) + \textit{recall}} \right)$$

$0 < \beta < 1 \Rightarrow$ More weight given to the precision

$\beta > 1 \Rightarrow$ More weight given to the recall

$\beta = 1 \Rightarrow$ Equal weight given to precision and recall (F_1)

OTHER MEASURES

Recall / Sensitivity /
True Positive Rate

		Predicted Output	
		Positive (PP)	Negative (PN)
True Values	Positive (P)	TP	FN
	Negative (N)	FP	TN

$$\frac{TP}{TP + FN}$$

Precision /
Positive Predictive Rate

		Predicted Output	
		Positive (PP)	Negative (PN)
True Values	Positive (P)	TP	FN
	Negative (N)	FP	TN

$$\frac{TP}{TP + FP}$$

$$\frac{TP}{TP + FP}$$

Negative Predictive Value

		Predicted Output	
		Positive (PP)	Negative (PN)
True Values	Positive (P)	TP	FN
	Negative (N)	FP	TN

$$\frac{TN}{TN + FN}$$

Specificity /
True Negative Rate

The diagram illustrates a confusion matrix for a binary classifier. It features a 2x2 grid of cells. The columns are labeled 'Predicted Output' with 'Positive (PP)' and 'Negative (PN)'. The rows are labeled 'True Values' with 'Positive (P)' and 'Negative (N)'. The cells contain the following labels: Top-Left (TP), Top-Right (FN), Bottom-Left (FP), and Bottom-Right (TN). The TP and TN cells are white with grey text, while the FP and FN cells are teal with white text. To the right of the matrix, the formula $\frac{TN}{TN + FP}$ is shown in orange, representing the Negative Predictive Value. A teal arrow points from the top-left towards the matrix.

	Predicted Output Positive (PP)	Predicted Output Negative (PN)
True Values Positive (P)	TP	FN
True Values Negative (N)	FP	TN

$\frac{TN}{TN + FP}$

$$\frac{TN}{TN + FP}$$


F_1

ROC Curve & AUC

ROC CURVE AND AUC

- A threshold converts a predicted probability or score into a class label.
 - Predict positive if $P(y_i = \text{positive}) > p$ ($0 < p < 1$)
 - ROC plots true positive rate vs. false positive rate as the threshold varies.
 - True positive rate = sensitivity. False positive rate = $1 - \text{specificity}$.
 - AUC (area under the ROC curve) summarizes performance.
 - AUC = 1: perfect model. AUC = 0.5: random performance.
-

EXAMPLE ROC CURVE



- **Each point uses a threshold.** Predict positive when the score exceeds it.
 - **Lower threshold:** More positives. TPR and FPR cannot decrease. Move up/right toward (1, 1).
 - **Higher threshold:** Fewer positives. TPR and FPR cannot increase. Move down/left toward (0, 0).
 - **Tradeoff:** Lower thresholds miss fewer positives but raise false alarms. Higher thresholds do the reverse.
-

METRICS IN SCIKIT-LEARN

FUNCTION NAMES

<https://scikit-learn.org/stable/api/sklearn.metrics.html#>

- `accuracy_score(y_true, y_pred)`
(default for many functions)
 - `classification_report(y_true, y_pred)`
 - `confusion_matrix(y_true, y_pred)`
 - `f1_score(y_true, y_pred)`
 - `precision_score(y_true, y_pred)`
 - `recall_score(y_true, y_pred)`
 - `roc_auc_score(y_true, y_score)` **Note:** (The second argument must be a score (e.g. probability) and not a class prediction)
-

`.PREDICT ()` VS `.PREDICT_PROBA ()`

- For classification estimators, there are two ways to produce predictions
 - `.predict ()` will produce a 1D array with n rows with the class prediction
 - Class with the highest probability
 - `.predict_proba ()` will produce a 2D array with n rows and n-class columns
 - Each column is the probability of the row observation being in class 0 to n-class
-

CLASSIFICATION REPORT

- This special function will show precision, recall, f1, and accuracy for all class (treat each class as positive class)

```
print(classification_report(y_test, yhat))
```

✓ 0.0s

	precision	recall	f1-score	support
0	0.95	0.99	0.97	980
1	0.91	1.00	0.95	1135
2	0.98	0.91	0.95	1032
3	0.96	0.96	0.96	1010
4	0.97	0.94	0.95	982
5	0.96	0.95	0.96	892
6	0.96	0.98	0.97	958
7	0.95	0.94	0.94	1028
8	0.99	0.91	0.95	974
9	0.93	0.95	0.94	1009
accuracy			0.95	10000
macro avg	0.96	0.95	0.95	10000
weighted avg	0.95	0.95	0.95	10000

MULTI-CLASS CLASSIFICATION

MULTI-CLASS STRATEGIES

- Some classifiers (e.g., Naive Bayes, KNN) are capable of handling multiple classes directly
 - Other classifiers (e.g., logistic regression) are binary classifiers
 - For binary classifiers, there are two strategies for extending to multiple classes:
 - One-versus-all (OvA) (sometimes called One-versus-rest (OvR))
 - One-versus-one (OvO)
-

ONE VERSUS ALL (OVA)

- For each class, train a model that classify that class versus all the rest:
 - Classify **1** versus **not 1**
 - Classify **2** versus **not 2**
 - etc.
- Classify a value into the classier whose classifier outputs the highest score (e.g., predictive probabilities)

ONE VERSUS ONE (OVO)

- Train a binary classifier for every pair of classes
 - 1 vs 2
 - 1 vs 3
 - 2 vs 3
 - 4 vs 5
 - etc
 - Classify value into the class that wins the most duels
-

OVA VERSUS OvO

- Suppose there are q classes
 - OvO requires training $q(q-1)/2$ classifiers
 - For the MNIST example this means $10(10-1)/2 = 45$ classifiers
 - However, each classifier needs only use a small subset of the data
 - In general OvA is the preferred strategy (only trains q classifiers)
 - But for classifiers that scale poorly with the size of the training set (e.g., SVM) OvO is preferred
-

SOFTMAX FUNCTION

- A mathematical function that converts a vector of real-valued scores into a probability distribution
- Used to convert real-valued outputs to “probabilities”
- Can be used as a way to extend logistic regression
- Often used in multi-class NN

The diagram shows the softmax function formula: $p(c_i) = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$. Annotations include: an arrow pointing to $p(c_i)$ with the text "Probability of class c_i "; an arrow pointing to e^{z_i} with the text " z_i is a raw score for the class (for example the class logit)"; and an arrow pointing to the denominator $\sum_{j=1}^K e^{z_j}$ with the text "Sum over all classes".

$$p(c_i) = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

Probability of class c_i

Sum over all classes

z_i is a raw score for the class
(for example the class logit)