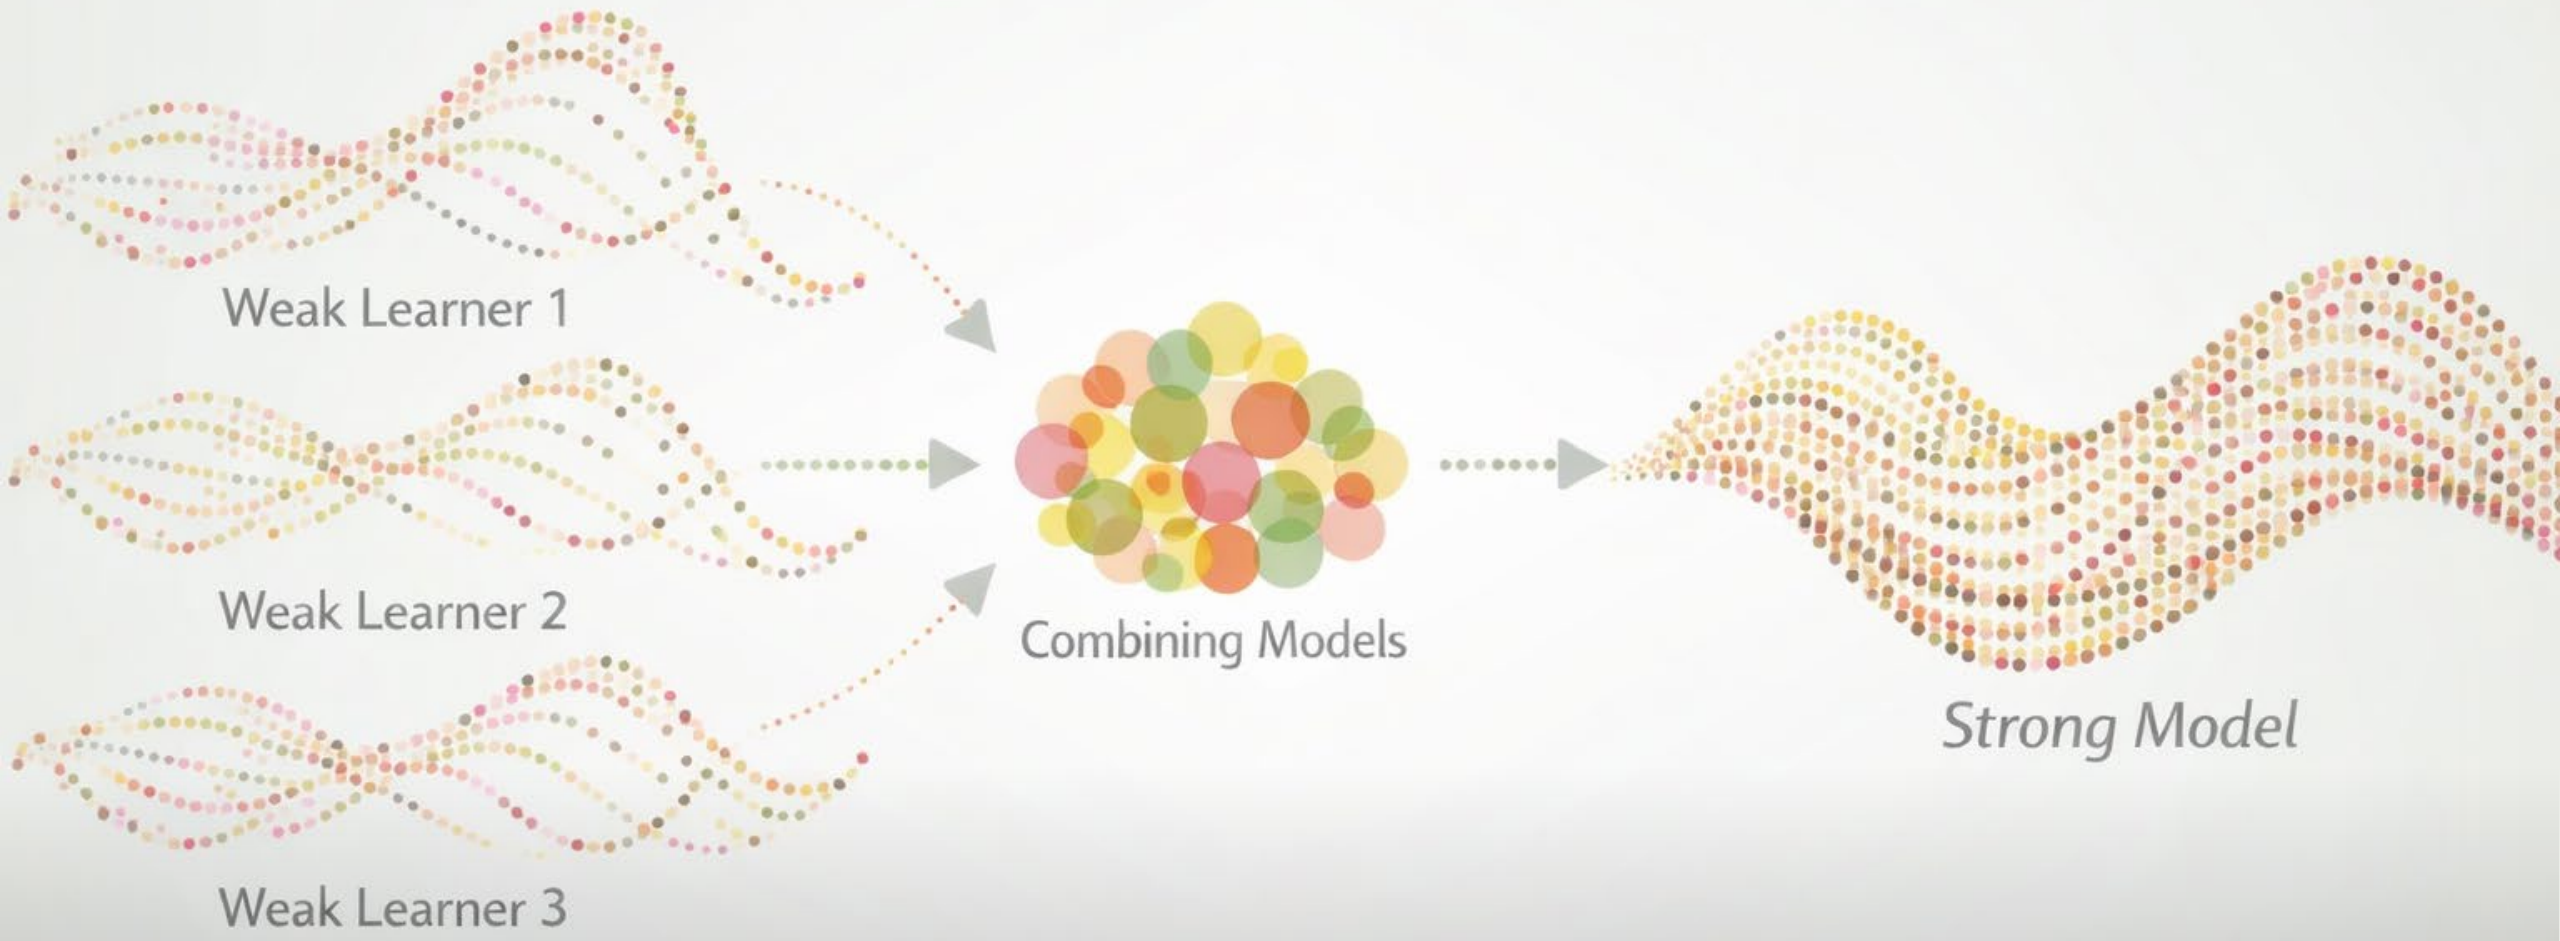

Weak Learner 1
Weak Learner 2
Weak Learner 3
Combining Models
Strong Model

BOOSTING

BOOSTING

- Boosting is a technique that was developed to improve **weak learners**
  - **Models** that are not much better than random guessing
  - Simple (few tuning parameters) and fast to train
  - Decision "stumps" are commonly used weak learners
- Idea: boosting methods train models sequentially and try to "***fix the mistakes***" of the preceding model

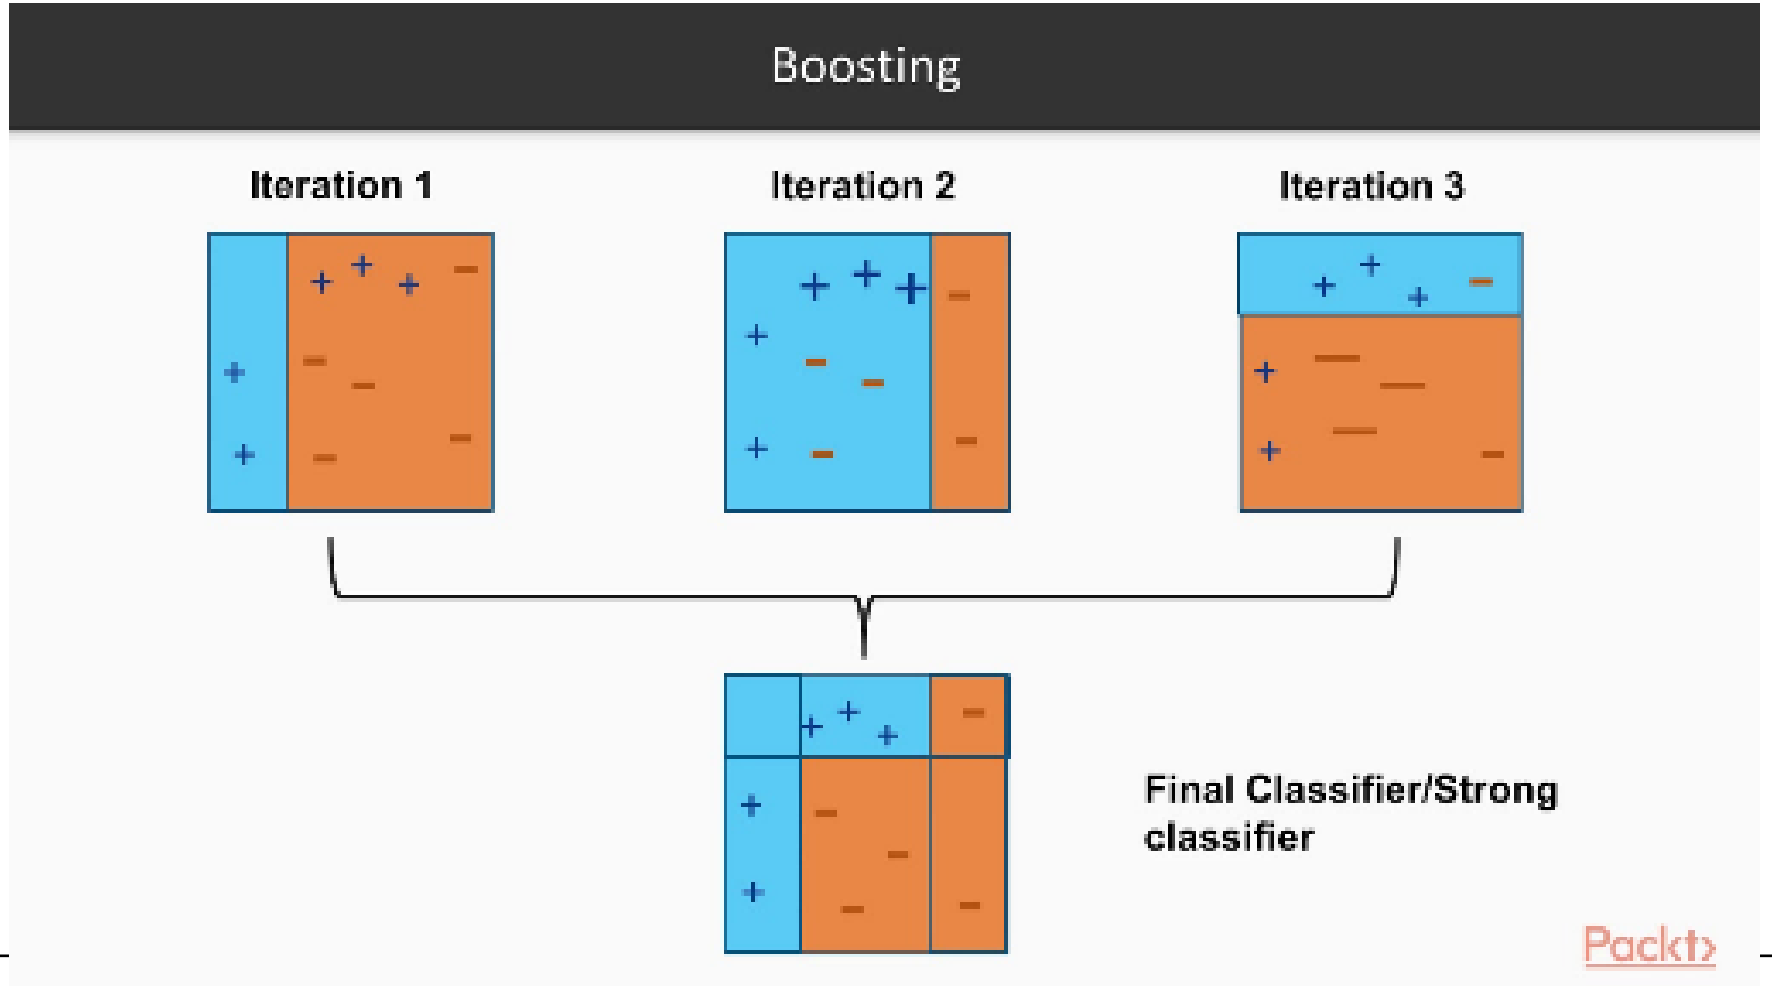
FOUNDATIONAL BOOSTING ALGORITHMS

- Adaptive Boosting (AdaBoost) — increase the weight of misclassified observations
- Gradient Boosting — learn and then correct the mistakes by predicting residuals
  - Modern implementations:
    - XGBoost (Extreme gradient boosting)
    - LightGBM (Light gradient boosting machine)
    - Catboost (Categorical boosting)

ADABOOST

- Adaptive Boosting (AdaBoost)
- Introduced in 1997 by Yoav Freund and Robert Schapire
- The main idea is to “pay more attention” to instances that were misclassified or had large errors in the previous model
- Focus learning on difficult instances
  - “Study” topics you missed last time

ADABOOST


Boosting
Iteration 1
Iteration 2
Iteration 3
Final Classifier/Strong classifier
Packt>

ADABOOST ALGORITHM

1. Assign each training point a weight: $w_i = 1/n$

   i indexes observations

2. Train a weak learner (e.g., a stump) and predict

3. Compute the weighted error rate: $\epsilon_j = \frac{\sum w_i I(y_i \neq \hat{y}_i)}{\sum w_i}$

   j indexes classifiers

4. Compute the classifier weight: $\alpha_j = \frac{1}{2}\log\left(\frac{1-\epsilon_j}{\epsilon_j}\right)$

   α_j is larger for more accurate models and sets their vote in the ensemble

ADABOOST

5. Update the instance weights, then normalize

   $w_{i,new} = w_{i,old} \cdot \exp(\alpha_j)$ if $y_i \neq \hat{y}_i$

   $w_{i,new} = w_{i,old} \cdot \exp(-\alpha_j)$ if $y_i = \hat{y}_i$

   Normalize by dividing by $\sum w_i$

6. The next classifier minimizes the weighted error, which now penalizes errors on high-weight instances

   The model focuses on instances misclassified in the previous step

ADABOOST

7. Update the final model prediction, scaled by a learning rate to regularize training

   $F_{\text{new}}(x) = F_{\text{old}}(x) + \lambda \cdot \alpha_j \cdot \hat{y}_j(x)$

   - A low learning rate ($\lambda < 1$) makes updates to the prediction smaller
   - This can reduce overfitting by making adjustments more gradual

ADABOOST SUMMARY

1. Initialize weights: $w_i = 1/n$
   i indexes observations, j indexes classifiers
2. Train a weak learner
3. Compute weighted error: $\epsilon_j = \frac{\sum w_i I(y_i \neq \hat{y}_i)}{\sum w_i}$
4. Compute classifier weight: $\alpha_j = \frac{1}{2}\log\left(\frac{1-\epsilon_j}{\epsilon_j}\right)$
5. Update and normalize weights:
   $w_{i,new} = w_{i,old} \cdot \exp(\alpha_j)$ if $y_i \neq \hat{y}_i$
   $w_{i,new} = w_{i,old} \cdot \exp(-\alpha_j)$ if $y_i = \hat{y}_i$
6. Repeat steps 2–5 with the new weights
7. Update the ensemble with a learning rate: $F_{\text{new}}(x) = F_{\text{old}}(x) + \lambda \cdot \alpha_j \cdot \hat{y}_j(x)$


AdaBoost Sequential Predictions & Instance Error Tracking (Noisy Two Moons)
Iteration 1 | Acc: 84.5% (31 errors)
Iteration 2 | Acc: 84.5% (31 errors)
Iteration 3 | Acc: 90.5% (19 errors)
Iteration 5 | Acc: 91.0% (18 errors)
Iteration 10 | Acc: 91.0% (18 errors)
Iteration 20 | Acc: 96.0% (8 errors)
Class 0 - Correct
Class 0 - Error
Class 1 - Correct
Class 1 - Error
Ensemble Boundary

GRADIENT BOOSTING

- Gradient Boosting directly targets the residuals (regression) or gradient of the loss function (classification)
- At each step, a new weak learner is fit to the residual or pseudo-residuals

GRADIENT BOOSTING

1.Train a weak learner and make predictions

2. Calculate the residuals (regression) as:

$$Residual = Actual\ Target - Current\ Prediction$$

3.Train a new weak learner that tries to predict the residuals from the previous step

4. Update the Predictions as:

$$\text{New Predictions} = \text{Current Prediction} + \lambda \times (\text{Predictions of New Model})$$

"Predicted Residuals"

Learning Rate

5. Repeat for a predefined number of iterations or until a stopping criterion is reached

6. Final prediction: $F_M(x) = F_0(x) + \sum_{m=1}^{M} \lambda h_m(x)$

EXAMPLE (WITH $\lambda = 1$)

Current Target	**Current Predictions**					
Y	Yhat-1	Resid-1	Yhat-2	Resid-2	Yhat-3	FinalPrediction
0.05	0.12	−0.07	−0.09	0.02	0.01	0.04
0.59	0.53	0.07	0.04	0.03	−0.07	0.50
0.17	0.12	0.04	0.04	0.00	0.01	0.17
−0.07	0.12	−0.19	−0.09	−0.10	0.01	0.04
0.34	0.12	0.22	0.16	0.06	0.01	0.29

EXAMPLE IN A TABLE

Current Target

Y	Yhat-1	**Resid-1**	Yhat-2	Resid-2	Yhat-3	FinalPrediction
0.05	0.12	**−0.07**	−0.09	0.02	0.01	0.04
0.59	0.53	**0.07**	0.04	0.03	−0.07	0.50
0.17	0.12	**0.04**	0.04	0.00	0.01	0.17
−0.07	0.12	**−0.19**	−0.09	−0.10	0.01	0.04
0.34	0.12	**0.22**	0.16	0.06	0.01	0.29

EXAMPLE IN A TABLE

		Current Target	Current Predictions			
Y	Yhat-1	Resid-1	Yhat-2	Resid-2	Yhat-3	FinalPrediction
0.05	0.12	−0.07	−0.09	0.02	0.01	0.04
0.59	0.53	0.07	0.04	0.03	−0.07	0.50
0.17	0.12	0.04	0.04	0.00	0.01	0.17
−0.07	0.12	−0.19	−0.09	−0.10	0.01	0.04
0.34	0.12	0.22	0.16	0.06	0.01	0.29

EXAMPLE IN A TABLE

Y	Yhat-1	Resid-1	Yhat-2	**Resid-2**	Yhat-3	FinalPrediction
0.05	0.12	−0.07	−0.09	**0.02**	0.01	0.04
0.59	0.53	0.07	0.04	**0.03**	−0.07	0.50
0.17	0.12	0.04	0.04	**0.00**	0.01	0.17
−0.07	0.12	−0.19	−0.09	**−0.10**	0.01	0.04
0.34	0.12	0.22	0.16	**0.06**	0.01	0.29

Current Target

EXAMPLE IN A TABLE

				Current Target	**Current Predictions**	
Y	Yhat-1	Resid-1	Yhat-2	Resid-2	Yhat-3	FinalPrediction
0.05	0.12	−0.07	−0.09	0.02	0.01	0.04
0.59	0.53	0.07	0.04	0.03	−0.07	0.50
0.17	0.12	0.04	0.04	0.00	0.01	0.17
−0.07	0.12	−0.19	−0.09	−0.10	0.01	0.04
0.34	0.12	0.22	0.16	0.06	0.01	0.29

EXAMPLE IN A TABLE

Y	Yhat-1	Resid-1	Yhat-2	Resid-2	Yhat-3	**Final Prediction**
0.05	0.12	−0.07	−0.09	0.02	0.01	**0.04**
0.59	0.53	0.07	0.04	0.03	−0.07	**0.50**
0.17	0.12	0.04	0.04	0.00	0.01	**0.17**
−0.07	0.12	−0.19	−0.09	−0.10	0.01	**0.04**
0.34	0.12	0.22	0.16	0.06	0.01	**0.29**

TUNING PARAMETERS

- **Number of estimators**
  - Too few → underfitting; too many → overfitting
  - Early stopping helps choose the number
- **Learning rate (λ)**
  - Shrinkage parameter that scales each tree's contribution
  - Smaller → more trees, better generalization
  - Larger → fewer trees, risk of overfitting or instability
- **Trade-off**
  - Low rate + many trees → slower, more stable
  - High rate + few trees → faster, more aggressive

GRADIENT BOOSTING FOR CLASSIFICATION

- The concept of "residuals" is a little different in the classification setting
  - The residual or pseudo residuals are the gradient of the loss function with respect to the predictions.
- For example, in binary classification:
  - Log loss or binary cross-entropy loss:

$$L(y, p) = -y\log(p) - (1 - y)\log(1 - p)$$

* y is the actual label (0 or 1)

* p is the predicted probability of the instance being in class 1

GRADIENT BOOSTING FOR CLASSIFICATION

- The gradient of the loss function with respect to p:

$$\frac{\partial L}{\partial p} = \frac{p - y}{p(1 - p)}$$

- This is simplified by working with the log odds,

  $z = \log\left(\frac{p}{1-p}\right)$, where $p = \frac{1}{1+e^{-z}}$

- The gradient with respect to the log odds is:

$$\frac{\partial L}{\partial z} = p - y$$

IMPLEMENTATIONS OF GRADIENT BOOSTING

- **Scikit-Learn**
  - GradientBoostingClassifier / GradientBoostingRegressor
  - Can be slower on large datasets
  - HistGradientBoosting provides a faster, histogram-based implementation
- **XGBoost (Extreme Gradient Boosting)**
  - Highly optimized and efficient
  - Supports L1 and L2 regularization
  - Handles missing values automatically
  - Parallelized and widely used in industry
- **LightGBM (Light Gradient Boosting Machine)**
  - Very fast for large datasets
  - Histogram-based learning for efficiency
  - Leaf-wise tree growth strategy
  - Native handling of categorical features
  - Memory efficient
- **CatBoost**
  - Strong performance with minimal tuning
  - Excellent handling of categorical variables
  - Reduces preprocessing requirements

INSTALLING BOOSTING LIBRARIES

Using pip

XGBoost

```
pip install xgboost
```

LightGBM

```
pip install lightgbm
```

CatBoost

```
pip install catboost
```

Using conda

XGBoost

```
conda install -c conda-forge py-xgboost
```

LightGBM

```
conda install -c conda-forge lightgbm
```

CatBoost

```
conda install -c conda-forge catboost
```